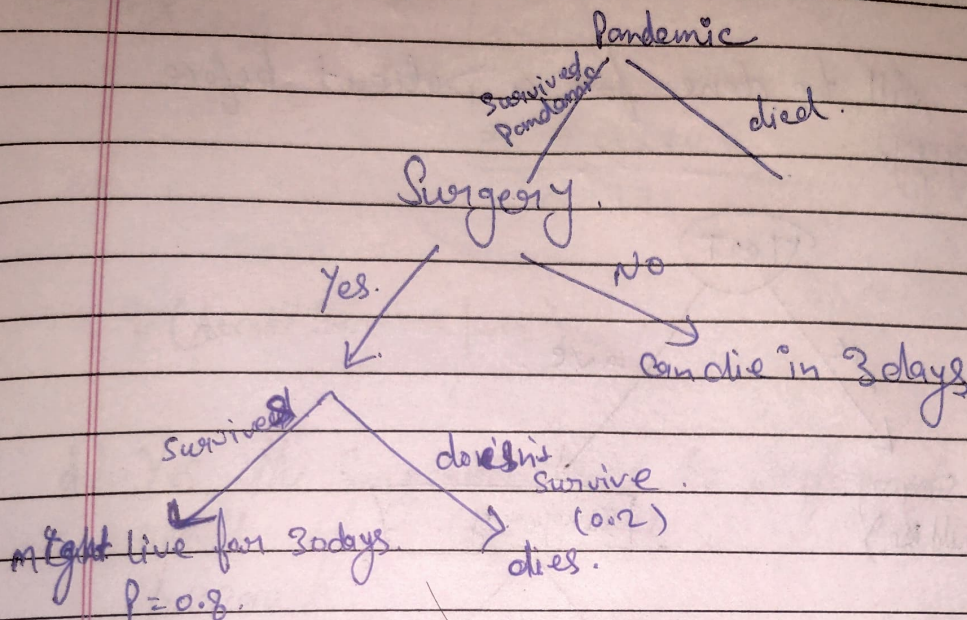


Assignment - 2

Sumit Kumar



$$P(\text{Survive after surgery}) = 0.8$$

$$P(\text{not survive after surgery}) = 1 - 0.8 = 0.2$$

2) living function. $L(n)$

$$L(30) = 1$$

$$L(n) = mx$$

$$m = \frac{1}{30}$$

$$L(n) = \frac{x}{30}$$

$$L(30) = \frac{30}{30} = 1$$

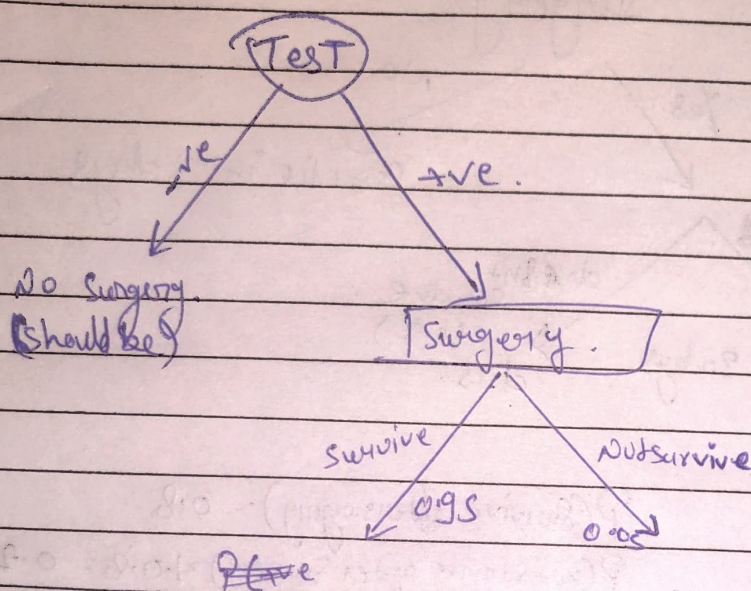
$$L(0) = \frac{0}{30} = 0$$

$$L(5) = ?$$

$$L(5) = \frac{5}{30} = \boxed{0.1}$$

c) A.T.Q.

Test will be done for a patient before Surgery.



$$P(\text{Survive Surgery given +ve}) = 0.95$$

$$P(\text{Not Survive Surgery given +ve}) =$$

$$P(\text{Survive Surgery given -ve}) = 0.05 \quad (1-0.95)$$

$$P(\text{Not survive given positive}) = 0.05$$

Find.

$$P(\text{Survive} | +ve) = \frac{P(+ve | \text{Survives}) P(\text{Survive Surgery})}{P(+ve)}$$

$$P(+ve) = P(+ve | \text{Survive}) + P(+ve | \text{Not Survive})$$

$$= 0.95 \times 0.8 + 0.05 \times 0.2$$

$$= 0.77$$

$$P(-ve) = 0.23$$

$$P(\text{Survive} | +ve) \\ P(\text{Survive Surgery} | +ve)$$

$$= \frac{0.95 \times 0.8}{0.77} = \frac{0.76}{0.77}$$

$$P(\text{doesn't survive} | +ve) = 0.987$$

d) Yes, the Surgery should be performed if the result of the test is positive because.

The true positive rate is high. means the outcome of test is true given some cases.

The false +ve Rate is too low.

⇒ The Surviving of the person is highly likely.

Contingency table.

TP FN = 1

FP ~~FN~~ = 1

1 1

0.95	0.05
0.05	0.95

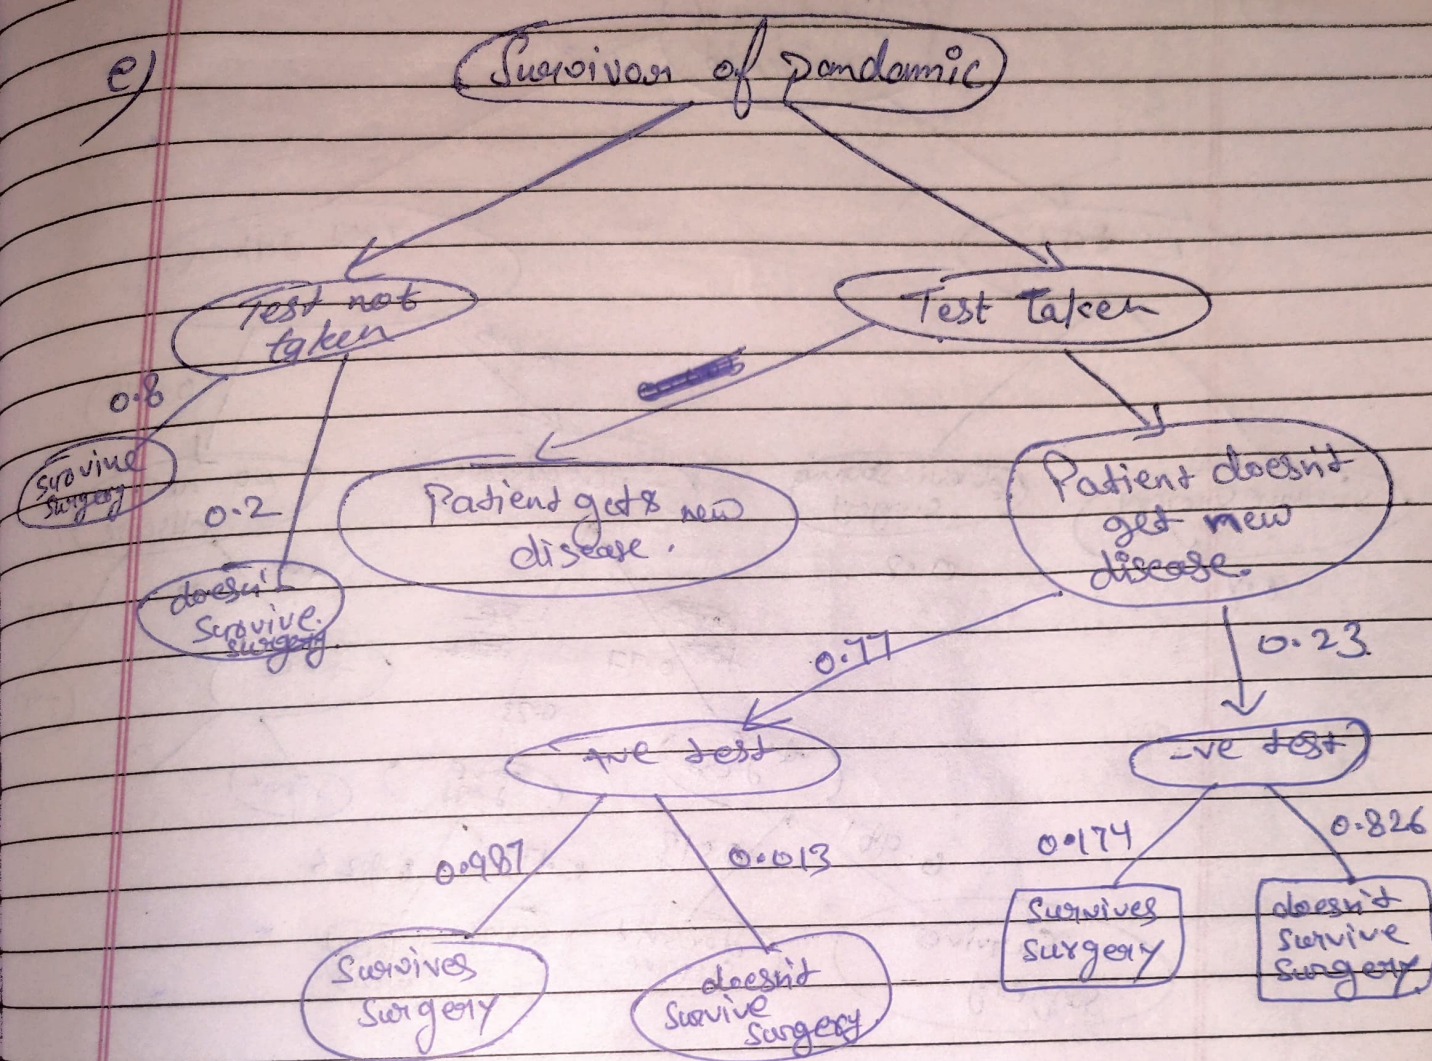
$$P(\text{Survives} | -ve) = \frac{P(-ve | \text{Survives}) P(\text{Survives})}{P(-ve)}$$

$$= \frac{0.05 \times 0.8}{0.23}$$

$$= 0.1739 \approx 0.174$$

$$P(\text{not Survive} | -ve) = 1 - 0.174$$
$$= 0.826$$

e)



f) Probability of new disease = 0.005 on test.

Should take the test or not.

$$P[\text{Survive}] = 0.987 \times 0.77 \times 0.995 + 0.0174 \times 0.23 \times 0.995$$

$$= 0.796 < 0.8$$

$$P[\text{Survive with no new disease after test taken}] < P[\text{Survive without test}]$$

⇒ Person should not take the test.

```
In [ ]: from utils import Dataset ,perceptron
import seaborn as sns
import matplotlib.pyplot as plt
import numpy as np
import pandas as pd

def plot1(inputs,weights):
    sns.scatterplot(data=inputs,x="X",y="Y",hue='Labels',s=2).set(title='plot with deci
inputs=np.array(inputs)
Ya_nn=[]
ia_nn=[]
for i in np.linspace(np.amin(inputs[:,1]),np.amax(inputs[:,1])):

    slope = -(weights[1])/(weights[2])
    intercept = -(weights[0]/weights[2])
    #y =mx+c, m is slope and c is intercept
    y = (slope*i) + intercept
    ia_nn.append(i)
    Ya_nn.append(y)
sns.lineplot(x=ia_nn,y=Ya_nn,color='black',markersize=4)
plt.legend()
plt.show()
```

```
In [ ]: def plot2(inputs,weights):
print('All the points on/above the decision boundary belongs to class 1')
sns.scatterplot(data=inputs,x="x1",y="x2",hue='out').set(title='plot with decision
inputs=np.array(inputs)
Ya_nn=[]
ia_nn=[]
for i in np.linspace(np.amin(inputs[:,1]),np.amax(inputs[:,1])):
    slope = 0
    intercept =0
    if(weights[2]!=0):
        slope = -(weights[1])/(weights[2])
    if(weights[2]!=0):
        intercept = -(weights[0]/weights[2])
    #y =mx+c, m is slope and c is intercept
    y = (slope*i) + intercept
    ia_nn.append(i)
    Ya_nn.append(y)
sns.lineplot(x=ia_nn,y=Ya_nn,color='black',markersize=4)
plt.legend()
plt.show()
```

Part 1

```
In [ ]: #without noise 10,000 samples
d=Dataset(10_000)
df=d.get()
df
```

```
Out[ ]:      X      Y  Labels
```

	X	Y	Labels
0	0.898230	-0.439526	0
1	0.723509	2.309685	1
2	-0.822943	-0.568124	0
3	-0.281687	3.959506	1
4	0.281912	0.959440	0
...
9995	0.867456	2.502486	1
9996	0.992129	3.125221	1
9997	-0.242525	0.970145	0
9998	0.135728	2.009254	1
9999	0.860023	2.489745	1

10000 rows × 3 columns

```
In [ ]: # d=Dataset(10000) with noise added
df1=d.get(True)
df1
```

```
Out [ ]:
```

	X	Y	Labels
0	0.561958	2.377508	1
1	-0.524846	0.868321	0
2	0.392301	0.718524	0
3	-0.027525	-1.074236	0
4	-0.127290	4.008247	1
...
9995	-0.457150	2.153987	1
9996	0.682129	3.870551	1
9997	-0.304008	-0.910514	0
9998	0.723689	-0.757239	0
9999	-0.120075	-0.982539	0

10000 rows × 3 columns

```
In [ ]: df.shape
```

```
Out [ ]: (10000, 3)
```




```
In [ ]: df1.shape
```


```
Out[ ]: (10000, 3)
```

Part 2

```
In [ ]: sns.scatterplot(data=df,x="X",y="Y",hue='Labels').set(title='Without noise')
plt.show()
```



```
In [ ]: sns.scatterplot(data=df1,x="X",y="Y",hue='Labels').set(title='With Noise')
plt.show()
```



Part 3

```
In [ ]: p=perceptron()
# training the perceptron with learnable bias on without noise data
p.train_weights(np.array(df),5,1)
```


```
>epoch=0, error=6.000 [-4.0, 0.720376711074459, 2.228126731090508]
```



```
>epoch=1, error=0.000 [-4.0, 0.720376711074459, 2.228126731090508]
>epoch=2, error=0.000 [-4.0, 0.720376711074459, 2.228126731090508]
>epoch=3, error=0.000 [-4.0, 0.720376711074459, 2.228126731090508]
>epoch=4, error=0.000 [-4.0, 0.720376711074459, 2.228126731090508]
```

In []:

```
plot1(df,p.weights)
```




In []:

```
p1=perceptron()
# the third parameter is a flag to make bias 0 or learnable
p1.train_weights(np.array(df1),5,1)
```

```
>epoch=0, error=13.000 [-5.0, -0.08963611587720852, 3.0705765156008695]
>epoch=1, error=0.000 [-5.0, -0.08963611587720852, 3.0705765156008695]
>epoch=2, error=0.000 [-5.0, -0.08963611587720852, 3.0705765156008695]
>epoch=3, error=0.000 [-5.0, -0.08963611587720852, 3.0705765156008695]
>epoch=4, error=0.000 [-5.0, -0.08963611587720852, 3.0705765156008695]
```

In []:

```
plot1(df1,p1.weights)
```



The Decision boundary for both the datasets with and without noise exist because the data was linearly separable even after adding noise.

Part 4

In []:

```
#classifier with bias set to constant 0 and data used = without noise data
p3=perceptron()
p3.train_weights(np.array(df),100,0)
```

```
>epoch=0, error=2966.000 [0.0, 0.7298420020772367, 0.357438305509001]
>epoch=1, error=2960.000 [0.0, 0.5832358101376358, 0.35825602366765097]
>epoch=2, error=2933.000 [0.0, 0.2057844099432271, 0.7686138975407416]
>epoch=3, error=2960.000 [0.0, 0.16148262079835995, 0.729842004188166]
>epoch=4, error=2969.000 [0.0, 0.5219670291163601, 0.33676161506953906]
>epoch=5, error=2986.000 [0.0, 0.4717218654541335, 0.704404884716293]
>epoch=6, error=2980.000 [0.0, 1.079825032968965, 0.9073725933009041]
>epoch=7, error=2976.000 [0.0, 0.2209896263871578, 0.7191236710494185]
>epoch=8, error=2969.000 [0.0, 0.43954112055514294, 0.3400822539368865]
>epoch=9, error=2962.000 [0.0, 1.015208100299604, 1.8071932011940732]
>epoch=10, error=3019.000 [0.0, 0.27536465020939405, 0.7640282261868007]
>epoch=11, error=2982.000 [0.0, 1.0515376689089853, 0.8808230634647499]
>epoch=12, error=2973.000 [0.0, 0.5526712300032313, 0.3441118716194903]
>epoch=13, error=2976.000 [0.0, 1.0479557213482988, 1.7845564166478538]
>epoch=14, error=2954.000 [0.0, 0.8895651222267846, 0.6777836946090053]
>epoch=15, error=2981.000 [0.0, 1.0238962330598933, 0.9260682819063845]
>epoch=16, error=2943.000 [0.0, 0.413556559985341, 0.27462716404239196]
>epoch=17, error=2962.000 [0.0, 0.8908658796434203, 0.7124451415946332]
>epoch=18, error=2931.000 [0.0, 0.23115779673824832, 0.740252831825445]
>epoch=19, error=2945.000 [0.0, 0.2235790460408258, 0.7242992172946875]
>epoch=20, error=2986.000 [0.0, 0.5234628793449905, 0.3871186492613027]
>epoch=21, error=2946.000 [0.0, 1.2920249004415536, 0.05901453359430808]
>epoch=22, error=2982.000 [0.0, 1.2204751962876135, 1.8716760371991037]
>epoch=23, error=2990.000 [0.0, 0.22422237420943558, 0.7234297371584014]
>epoch=24, error=2986.000 [0.0, 0.5241062075136003, 0.3862491691250166]
>epoch=25, error=2970.000 [0.0, 0.44639381683861723, 0.678922103955324]
>epoch=26, error=2970.000 [0.0, 0.26706168043483247, 0.7502170105397071]
>epoch=27, error=2977.000 [0.0, 0.21916295331120894, 0.6911493922734425]
>epoch=28, error=2930.000 [0.0, 0.21757398061945143, 0.7685420138586131]
>epoch=29, error=2968.000 [0.0, 1.0142076828134532, 1.8075247419397082]
>epoch=30, error=3019.000 [0.0, 0.2743642327232432, 0.7643597669324357]
>epoch=31, error=2952.000 [0.0, 1.1771151204575394, 1.8575551953991236]
>epoch=32, error=2946.000 [0.0, 0.36000097747622983, 0.29509140122480215]
>epoch=33, error=2964.000 [0.0, 0.9634727157056591, 0.09376357206955566]
>epoch=34, error=2978.000 [0.0, 0.8286671960717851, 0.17091159142503753]
>epoch=35, error=2978.000 [0.0, 1.0807771201850676, 0.9121824437372951]
>epoch=36, error=2969.000 [0.0, 0.41333223013522646, 0.2639591861289309]
>epoch=37, error=2937.000 [0.0, 1.0290007708943756, 1.8365090267355781]
>epoch=38, error=2969.000 [0.0, 0.4234426902239492, 0.2579431579500425]
>epoch=39, error=2936.000 [0.0, 1.0984844252118817, 0.9248185494052407]
>epoch=40, error=2938.000 [0.0, 0.40318690552436887, 1.015195217376354]
>epoch=41, error=2961.000 [0.0, 0.2249056046032074, 0.7175465413194148]
>epoch=42, error=2966.000 [0.0, 0.48820509362789766, 0.3726391945716837]
>epoch=43, error=2983.000 [0.0, 1.070524616034751, 0.14059532518057638]
>epoch=44, error=2992.000 [0.0, 1.1933755799043713, 1.8207417640622467]
>epoch=45, error=2980.000 [0.0, 0.8353787811192361, 0.6998728606801979]
>epoch=46, error=2957.000 [0.0, 0.5709799817256336, 0.36623613088986684]
>epoch=47, error=2934.000 [0.0, 0.7853815210235524, 0.11009130939192546]
>epoch=48, error=2984.000 [0.0, 0.7710810197975861, 0.6625578906055828]
>epoch=49, error=2983.000 [0.0, 1.0559619710752766, 1.7878318840814518]
>epoch=50, error=2985.000 [0.0, 0.38474941336508084, 0.3022871081049666]
>epoch=51, error=2949.000 [0.0, 1.306959398099452, 0.33111040116575696]
>epoch=52, error=2942.000 [0.0, 0.38228576669027503, 0.3017813355892739]
```



```

>epoch=53, error=2971.000 [0.0, 1.013526276108336, 0.01271917072614348]
>epoch=54, error=2973.000 [0.0, 0.6237100943362186, 0.3351875401664276]
>epoch=55, error=2959.000 [0.0, 0.2074676747499653, 0.7705033767694777]
>epoch=56, error=2955.000 [0.0, 0.8055852329601458, 0.6899159179378979]
>epoch=57, error=2949.000 [0.0, 0.45229981165421385, 0.71146885373078]
>epoch=58, error=2977.000 [0.0, 0.14208576312262844, 0.704971490343748]
>epoch=59, error=2965.000 [0.0, 0.4461986841854164, 0.9940558255217292]
>epoch=60, error=2939.000 [0.0, 0.5633765427537001, 0.35474479368297207]
>epoch=61, error=2987.000 [0.0, 0.5119312457203402, 0.394646613895442]
>epoch=62, error=2970.000 [0.0, 0.534745362493044, 0.3779592823018635]
>epoch=63, error=2947.000 [0.0, 1.0273500674541456, 0.9786792570226949]
>epoch=64, error=2951.000 [0.0, 1.0892495592794464, 0.922387617734958]
>epoch=65, error=2960.000 [0.0, 0.2162171657523444, 0.7701594051693813]
>epoch=66, error=2974.000 [0.0, 1.2346106013652909, 1.861068228691483]
>epoch=67, error=2987.000 [0.0, 0.5234085199478451, 0.38316788394450196]
>epoch=68, error=2947.000 [0.0, 0.5312481918168195, 0.3781791364588837]
>epoch=69, error=2957.000 [0.0, 0.18670376451299986, 0.715567940751555]
>epoch=70, error=2991.000 [0.0, 1.0692349887912083, 1.8215128855644127]
>epoch=71, error=2977.000 [0.0, 0.5484383132575446, 0.36547457444020026]
>epoch=72, error=2963.000 [0.0, 0.5307174625607014, 0.37668383309745856]
>epoch=73, error=2929.000 [0.0, 1.0888987756027788, 0.914777160787471]
>epoch=74, error=2959.000 [0.0, 0.4892079182042428, 0.3577475901042183]
>epoch=75, error=2980.000 [0.0, 0.3516770827282829, 0.7137366194037883]
>epoch=76, error=2952.000 [0.0, 0.5167642367549385, 0.32442246994427004]
>epoch=77, error=2995.000 [0.0, 0.3905955395355807, 0.30001945257603435]
>epoch=78, error=2963.000 [0.0, 1.2341992048468862, 0.21152734632781878]
>epoch=79, error=2978.000 [0.0, 1.357478140685172, 0.3222387287233023]
>epoch=80, error=2992.000 [0.0, 0.5982498622677015, 0.36817567160504894]
>epoch=81, error=2968.000 [0.0, 0.8630263116400407, 0.6791626659175686]
>epoch=82, error=2946.000 [0.0, 0.21564619698009468, 0.7665582619873074]
>epoch=83, error=2976.000 [0.0, 0.2254605698586607, 0.722054598538863]
>epoch=84, error=2929.000 [0.0, 0.43899393925054686, 1.0149527097663333]
>epoch=85, error=2961.000 [0.0, 1.094251549779135, 0.9230244344393049]
>epoch=86, error=2949.000 [0.0, -0.2327600096165423, 0.33110459579379425]
>epoch=87, error=2990.000 [0.0, 0.5270926753150253, 0.3836383258418171]
>epoch=88, error=2954.000 [0.0, 0.3945932807696495, 0.30412907566079117]
>epoch=89, error=2955.000 [0.0, 1.6046741234935753, 1.3940772544985056]
>epoch=90, error=2942.000 [0.0, 1.0546074237744407, 1.8206648641931085]
>epoch=91, error=2938.000 [0.0, 1.0717646100135783, 0.9109388556224572]
>epoch=92, error=2995.000 [0.0, 1.037873243028897, 1.8721538680239593]
>epoch=93, error=2976.000 [0.0, 0.45747633718115455, 0.9993320340648483]
>epoch=94, error=2932.000 [0.0, 1.0259644472285208, 1.800666461210835]
>epoch=95, error=2927.000 [0.0, 0.4048775487847811, 0.7628896522114547]
>epoch=96, error=2933.000 [0.0, 0.5754382951548034, 0.3551700715093872]
>epoch=97, error=2969.000 [0.0, 0.5103494190500921, 0.3519204291583412]
>epoch=98, error=2993.000 [0.0, 0.5474508945226755, 0.3843436913677226]
>epoch=99, error=2971.000 [0.0, 0.5483186593984457, 0.3843399300705964]

```

The decision boundary for this model won't exist because when we made the bias 0 we are forcing the decision boundary to pass through the origin and also separate the data, but the point (0,0) is the center for the circle having $k=0$ (label = 0) which makes the weights oscillate to the error value no matter the number of epochs which results in a non-existent decision boundary which passes through (0,0)

However this is not the case when we allow the bias to be learnable (can update with error) it can shift with the weights accordingly to finally converge to a proper decision boundary.

When bias was learnable it took only 2 epochs to find the correct weights but when we fix the bias the error is not reducing even after 100 epochs hence the perceptron will not converge if the bias is set to 0 for the given dataset


The plots can be seen below for both the cases.

In []:


```
print('Forcing the boundary line to pass through origin')
```

```
plot1(df,p3.weights)
print('Bias can be adjusted with no restriction')
plot1(df,p.weights)
```

Forcing the boundary line to pass through origin



Bias can be adjusted with no restriction



Part 5

```
In [ ]: xor_d={'x1':[0,0,1,1], 'x2':[0,1,0,1], 'out':[0,1,1,0]}
and_d={'x1':[0,0,1,1], 'x2':[0,1,0,1], 'out':[0,0,0,1]}
or_d={'x1':[0,0,1,1], 'x2':[0,1,0,1], 'out':[0,1,1,1]}
```

```
In [ ]: and_df=pd.DataFrame(and_d)
print(and_df)
p_and=perceptron()
p_and.train_weights(np.array(and_df),6,1)
print('With learnable bias')
plot2(and_df,p_and.weights)

p_and1=perceptron()
p_and1.train_weights(np.array(and_df),6,0)
```




```
print('With 0 bias')
plot2(and_df,p_and1.weights)
```

	x1	x2	out
0	0	0	0
1	0	1	0
2	1	0	0
3	1	1	1

```
>epoch=0, error=3.000 [-1.0, 2.0, 1.0]
>epoch=1, error=3.000 [-2.0, 2.0, 1.0]
>epoch=2, error=2.000 [-2.0, 2.0, 2.0]
>epoch=3, error=1.000 [-3.0, 2.0, 1.0]
>epoch=4, error=0.000 [-3.0, 2.0, 1.0]
>epoch=5, error=0.000 [-3.0, 2.0, 1.0]
```

With learnable bias

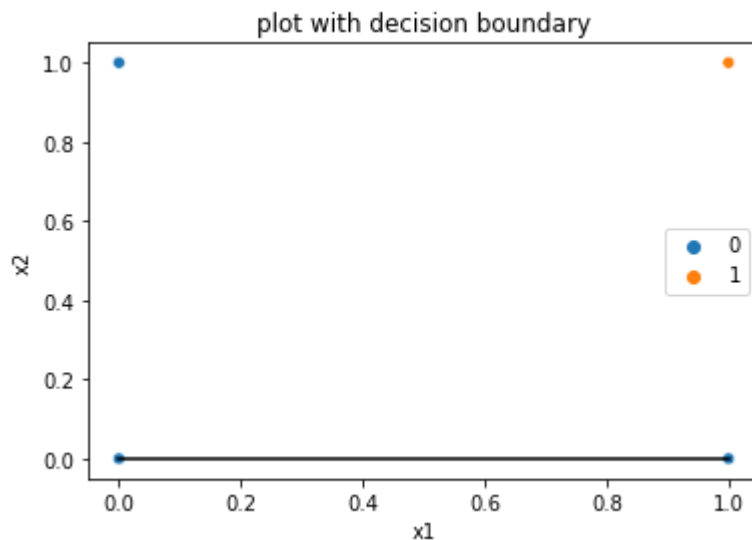
All the points on/above the decision boundary belongs to class 1



```
>epoch=0, error=3.000 [0.0, 0.0, 0.0]
>epoch=1, error=4.000 [0.0, 0.0, 0.0]
>epoch=2, error=4.000 [0.0, 0.0, 0.0]
>epoch=3, error=4.000 [0.0, 0.0, 0.0]
>epoch=4, error=4.000 [0.0, 0.0, 0.0]
>epoch=5, error=4.000 [0.0, 0.0, 0.0]
```

With 0 bias

All the points on/above the decision boundary belongs to class 1




```
In [ ]: or_df=pd.DataFrame(or_d)
print(or_df)
p_or=perceptron()
#Learnable bias
p_or.train_weights(np.array(or_df),6,1)
plot2(or_df,p_or.weights)
p_or1=perceptron()
p_or1.train_weights(np.array(or_df),6,0)
plot2(or_df,p_or1.weights)
```

```
   x1  x2  out
0    0    0    0
1    0    1    1
2    1    0    1
3    1    1    1
```


```
>epoch=0, error=1.000 [-1.0, 1.0, 1.0]
>epoch=1, error=0.000 [-1.0, 1.0, 1.0]
>epoch=2, error=0.000 [-1.0, 1.0, 1.0]
>epoch=3, error=0.000 [-1.0, 1.0, 1.0]
>epoch=4, error=0.000 [-1.0, 1.0, 1.0]
>epoch=5, error=0.000 [-1.0, 1.0, 1.0]
```

All the points on/above the decision boundary belongs to class 1



```
>epoch=0, error=1.000 [0.0, 1.0, 1.0]
>epoch=1, error=1.000 [0.0, 1.0, 1.0]
>epoch=2, error=1.000 [0.0, 1.0, 1.0]
>epoch=3, error=1.000 [0.0, 1.0, 1.0]
>epoch=4, error=1.000 [0.0, 1.0, 1.0]
>epoch=5, error=1.000 [0.0, 1.0, 1.0]
```


All the points on/above the decision boundary belongs to class 1



```
In [ ]: xor_df=pd.DataFrame(xor_d)
print(xor_df)
p_xor=perceptron()
p_xor.train_weights(np.array(xor_df),5,1)
plot2(xor_df,p_xor.weights)
p_xor1=perceptron()
p_xor1.train_weights(np.array(xor_df),5,0)
plot2(xor_df,p_xor1.weights)
```

	x1	x2	out
0	0	0	0
1	0	1	1
2	1	0	1
3	1	1	0

```
>epoch=0, error=2.000 [-2.0, 0.0, 0.0]
>epoch=1, error=3.000 [-1.0, 0.0, 0.0]
>epoch=2, error=2.000 [-1.0, -1.0, 0.0]
>epoch=3, error=3.000 [0.0, -1.0, 0.0]
>epoch=4, error=4.000 [0.0, -1.0, 0.0]
All the points on/above the decision boundary belongs to class 1
```




```
>epoch=0, error=2.000 [0.0, 0.0, 0.0]
>epoch=1, error=2.000 [0.0, -1.0, -1.0]
>epoch=2, error=4.000 [0.0, -1.0, -1.0]
```

```
>epoch=3, error=4.000 [0.0, -1.0, -1.0]
```

```
>epoch=4, error=4.000 [0.0, -1.0, -1.0]
```

All the points on/above the decision boundary belongs to class 1



Wrong decision boundary for XOR dataset because data can not be separated with one decision boundary

part 6

Given a hyperplane boundary and a point we can classify the point into class 0 or 1 by putting the point coordinates into the hyperplane equation by applying the sign(signum) function on the result(let's call it R)

Assumption :: sign function gives 1 when R is ≥ 0 and 0 when $R < 0$

when the sign function gives 1 the class of the point is 1 else the class of the point is 0.

In []:


```
In [ ]: import pandas as pd
import seaborn as sb
import numpy as np
import matplotlib.pyplot as plt
import scipy as stats
from sklearn import preprocessing
from sklearn.metrics import classification_report
from sklearn.metrics import confusion_matrix
from sklearn.tree import DecisionTreeClassifier # Import Decision Tree Classifier
from sklearn import metrics #Import scikit-learn metrics module for accuracy calculation
```

```
In [ ]: df = pd.read_csv('BitcoinHeistData.csv')
```

```
In [ ]: df
```

```
Out[ ]:
```

	address	year	day	length	weight	count	looped	neigh
0	111K8kZAEJg245r2cM6y9zgjGHZtJPY6	2017	11	18	0.008333	1	0	
1	1123pJv8jzeFQaCV4w644pzQJzVWay2zcA	2016	132	44	0.000244	1	0	
2	112536im7hy6wtKbpH1qYDWtTyMRACa2p7	2016	246	0	1.000000	1	0	
3	1126eDRw2wqSkWosjTCre8cjjQW8sSeWH7	2016	322	72	0.003906	1	0	
4	1129TSjKtx65E35GiUo4AYVeyo48twbrGX	2016	238	144	0.072848	456	0	
...
2916692	12D3trgho1vJ4mGtWBRPyHdMJK96TRYSry	2018	330	0	0.111111	1	0	
2916693	1P7PputTcVkhXBmXBvSD9MJ3UYPsiou1u2	2018	330	0	1.000000	1	0	
2916694	1KYiKJEfdJtap9QX2v9BXJMpZ2SfU4pgZw	2018	330	2	12.000000	6	6	
2916695	15iPUJsRNZQZHmZZVwmQ63srsughCXV4a	2018	330	0	0.500000	1	0	
2916696	3LFFBxp15h9KSftaw55np8eP5fv6kdK17e	2018	330	144	0.073972	6800	0	

2916697 rows × 10 columns



```
In [ ]: # df.info()
df.isna().sum().sum()
```

```
Out[ ]: 0
```

No nan values found

```
In [ ]: #still for precaution
df.dropna()
```

```
Out[ ]:
```

	address	year	day	length	weight	count	looped	neigh
--	---------	------	-----	--------	--------	-------	--------	-------

	address	year	day	length	weight	count	looped	neigh
0	111K8kZAEJg245r2cM6y9zgJGHZtJPy6	2017	11	18	0.008333	1	0	
1	1123pJv8jzeFQaCV4w644pzQJzVWay2zcA	2016	132	44	0.000244	1	0	
2	112536im7hy6wtKbpH1qYDWtTyMRAcA2p7	2016	246	0	1.000000	1	0	
3	1126eDRw2wqSkWosjTCre8cjjQW8sSeWH7	2016	322	72	0.003906	1	0	
4	1129TSjKtx65E35GiUo4AYVeyo48twbrGX	2016	238	144	0.072848	456	0	
...
2916692	12D3trgho1vJ4mGtWBRPyHdMJK96TRYSry	2018	330	0	0.111111	1	0	
2916693	1P7PputTcVkhXBmXBvSD9MJ3UYPsiou1u2	2018	330	0	1.000000	1	0	
2916694	1KYiKJEfdJtap9QX2v9BXJMpz2SfU4pgZw	2018	330	2	12.000000	6	6	
2916695	15iPUJsRNZQZHmZZVwmQ63srsughCXV4a	2018	330	0	0.500000	1	0	
2916696	3LFFBxp15h9KSFtaw55np8eP5fv6kdK17e	2018	330	144	0.073972	6800	0	

2916697 rows × 10 columns



```
In [ ]: n = len(pd.unique(df['address']))
n
```

Out []: 2631095

```
In [ ]: df=df.drop(['address'],axis=1)
#dropped address column because two many unique values will only
#contribute to computation complexity when the address feature is not significant
df
```

	year	day	length	weight	count	looped	neighbors	income	label
0	2017	11	18	0.008333	1	0	2	1.000500e+08	princetonCerber
1	2016	132	44	0.000244	1	0	1	1.000000e+08	princetonLocky
2	2016	246	0	1.000000	1	0	2	2.000000e+08	princetonCerber
3	2016	322	72	0.003906	1	0	2	7.120000e+07	princetonCerber
4	2016	238	144	0.072848	456	0	1	2.000000e+08	princetonLocky
...
2916692	2018	330	0	0.111111	1	0	1	1.255809e+09	white
2916693	2018	330	0	1.000000	1	0	1	4.409699e+07	white
2916694	2018	330	2	12.000000	6	6	35	2.398267e+09	white
2916695	2018	330	0	0.500000	1	0	1	1.780427e+08	white
2916696	2018	330	144	0.073972	6800	0	2	1.123500e+08	white

2916697 rows × 9 columns

Part 1

```
In [ ]: target=['label']
features=df.columns[:-1]
features
```

```
Out[ ]: Index(['year', 'day', 'length', 'weight', 'count', 'looped', 'neighbors',
        'income'],
        dtype='object')
```

```
In [ ]: X=df[features]
print(X)
Y=df[target]
print(Y)
```

	year	day	length	weight	count	looped	neighbors	income
0	2017	11	18	0.008333	1	0	2	1.000500e+08
1	2016	132	44	0.000244	1	0	1	1.000000e+08
2	2016	246	0	1.000000	1	0	2	2.000000e+08
3	2016	322	72	0.003906	1	0	2	7.120000e+07
4	2016	238	144	0.072848	456	0	1	2.000000e+08
...
2916692	2018	330	0	0.111111	1	0	1	1.255809e+09
2916693	2018	330	0	1.000000	1	0	1	4.409699e+07
2916694	2018	330	2	12.000000	6	6	35	2.398267e+09
2916695	2018	330	0	0.500000	1	0	1	1.780427e+08
2916696	2018	330	144	0.073972	6800	0	2	1.123500e+08

[2916697 rows x 8 columns]

	label
0	princetonCerber
1	princetonLocky
2	princetonCerber
3	princetonCerber
4	princetonLocky
...	...
2916692	white
2916693	white
2916694	white
2916695	white
2916696	white

[2916697 rows x 1 columns]

```
In [ ]: #encoding the target column
label_encoder = preprocessing.LabelEncoder()
df['label']= label_encoder.fit_transform(df['label'])
```

```
In [ ]: df = df.sample(frac=1)
train_size = 0.70
test_size = 0.15
valid_size=0.15
```

```

train_index = int(len(df)*train_size)

df_train = df[0:train_index]
df_rem = df[train_index:]

valid_index = int(len(df)*valid_size)

df_valid = df[train_index:train_index+valid_index]
df_test = df[train_index+valid_index:]

X_train, y_train = df_train.drop(columns='label'), df_train['label']
X_valid, y_valid = df_valid.drop(columns='label'), df_valid['label']
X_test, y_test = df_test.drop(columns='label'), df_test['label']
trees_both=[]

```

```

In [ ]: # Create Decision Tree classifier object with entropy
accuracy_entropy=[]
depth=[4,8,10,15,20]
for i in depth:
    clf = DecisionTreeClassifier(criterion="entropy", max_depth=i)
    # Train Decision Tree Classifier
    clf = clf.fit(X_train,y_train)
    #Predict the response for test dataset
    y_pred = clf.predict(X_valid)
    print("Accuracy for depth",i,":",metrics.accuracy_score(y_valid, y_pred))
    accuracy_entropy.append(metrics.accuracy_score(y_valid, y_pred))

```

```

Accuracy for depth 4 : 0.9855155610005851
Accuracy for depth 8 : 0.9857807014335869
Accuracy for depth 10 : 0.9870424041837331
Accuracy for depth 15 : 0.9879201104447045
Accuracy for depth 20 : 0.9860984128145114

```

```

In [ ]: # Create Decision Tree classifier object with gini index
accuracy_gini=[]
for i in depth:
    clf = DecisionTreeClassifier(criterion="gini", max_depth=i)
    # Train Decision Tree Classifier
    clf = clf.fit(X_train,y_train)
    #Predict the response for test dataset
    y_pred = clf.predict(X_valid)
    print("Accuracy for depth",i,":",metrics.accuracy_score(y_valid, y_pred))
    accuracy_gini.append(metrics.accuracy_score(y_valid, y_pred))

```

```

Accuracy for depth 4 : 0.9855155610005851
Accuracy for depth 8 : 0.9862081260971328
Accuracy for depth 10 : 0.9867864065242832
Accuracy for depth 15 : 0.9878286827091867
Accuracy for depth 20 : 0.9863589818607372

```

```

In [ ]: for i in range(len(depth)):
    print("Accuracy for depth",depth[i],"with entropy:",accuracy_entropy[i])
    print("Accuracy for depth",depth[i],"with gini:",accuracy_gini[i])

```

```

Accuracy for depth 4 with entropy: 0.9855155610005851
Accuracy for depth 4 with gini: 0.9855155610005851
Accuracy for depth 8 with entropy: 0.9857807014335869

```

Accuracy for depth 8 with gini: 0.9862081260971328
 Accuracy for depth 10 with entropy: 0.9870424041837331
 Accuracy for depth 10 with gini: 0.9867864065242832
 Accuracy for depth 15 with entropy: 0.9879201104447045
 Accuracy for depth 15 with gini: 0.9878286827091867
 Accuracy for depth 20 with entropy: 0.9860984128145114
 Accuracy for depth 20 with gini: 0.9863589818607372

The best accuracy for the ginni and entropy creteria is best with the max depth of 15

```
In [ ]:
clf = DecisionTreeClassifier(criterion="entropy", max_depth=15)
# Train Decision Tree Classifier
clf = clf.fit(X_train,y_train)
#Predict the response for test dataset
y_pred = clf.predict(X_valid)
print("Accuracy for depth",15,":",metrics.accuracy_score(y_valid, y_pred))
accuracy_entropy.append(metrics.accuracy_score(y_valid, y_pred))
```

Accuracy for depth 15 : 0.9878766822703335

```
In [ ]:
clf = DecisionTreeClassifier(criterion="gini", max_depth=15)
# Train Decision Tree Classifier
clf = clf.fit(X_train,y_train)
#Predict the response for test dataset
y_pred = clf.predict(X_test)
print("Accuracy for depth",15,":",metrics.accuracy_score(y_test, y_pred))
# accuracy_gini.append(metrics.accuracy_score(y_valid, y_pred))
```

Accuracy for depth 15 : 0.9881921619360649

Part 2

Ensembling is a method to combine multiple not-so-good models to get a better performing model. Create 100 different decision stumps (max depth 3). For each stump, train it on randomly selected 50% of the training data, i.e., select data for each stump separately. Now, predict the test samples' labels by taking a majority vote of the output of the stumps. How is the performance affected as compared to parts (a)

```
In [ ]:
def train_trees(X, Y, num_trees):
    indices = [i for i in range(X.shape[0])]

    # Here, note that we have set max_depth = 3 which
    # makes all the classifiers weak
    trees = [DecisionTreeClassifier(max_depth=3) for _ in range(num_trees)]
    for tree in trees:
        # Selecting n random samples with replacement from training set
        random_indices = np.random.choice(indices, X.shape[0])
        print(random_indices)
        # Bootstrap training data
        X_bootstrap= X.iloc[random_indices]
        Y_bootstrap= Y.iloc[random_indices]
        # X_bootstrap = X[random_indices]
        # Y_bootstrap = Y[random_indices]

        tree.fit(X_bootstrap, Y_bootstrap)
```



```
return trees
```

```
In [ ]: def predict(X, trees):
    predictions = []
    for tree in trees:
        Y_pred = tree.predict(X)
        predictions.append(Y_pred)

    predictions = np.array(predictions)

    # Aggregating all predictions to get final prediction
    # Since this is a classification problem, we use mode
    # i.e. the prediction that occurs the maximum number
    # of times. In case of regression problem, we use mean
    prediction = np.array(stats.stats.mode(predictions))
    return prediction[0, 0, :]
```

```
In [ ]: # NUM_RANDOM_FEATURES =
# put 50% of data from the dataset df into x train and y train
df = df.sample(frac=1)
train_size = 0.50
# test_size = 0.50

train_index = int(len(df)*train_size)

df_train = df[0:train_index]
df_rem = df[train_index:]

# df_test = df[train_index+valid_index:]

X_train, y_train = df_train.drop(columns='label'), df_train['label']
# X_valid, y_valid = df_valid.drop(columns='label'), df_valid['label']
X_test, y_test = df_rem.drop(columns='label'), df_rem['label']
```

```
In [ ]: # accuracies = []
trees = train_trees(X_train, y_train, 100)
prediction = predict(X_test, trees)
# accuracy = (prediction == y_test).sum() / prediction.shape[0]
# accuracies.append(accuracy)

# plt.plot(accuracies)
```

```
[ 487038 1007506 563817 ... 1368765 1150143 97461]
[1407639 39934 202282 ... 92420 423861 476474]
[ 790485 1354128 382393 ... 244555 422526 909608]
[ 300892 159636 801999 ... 1300258 1271287 1307805]
[ 812440 1316672 967794 ... 1256109 171250 1216143]
[380163 598884 413676 ... 458473 324453 856961]
[ 519269 548718 1156057 ... 183524 86952 243393]
[507258 233889 106903 ... 805654 701789 505699]
[ 546497 658392 698508 ... 786156 1368889 217902]
[1244210 739834 1293874 ... 594322 737519 569317]
[1038793 910238 966858 ... 81740 819588 153268]
[1240813 1078517 307703 ... 938405 924306 337887]
[492803 674715 517661 ... 408470 492369 776330]
```

[664939 735775 963678 ... 8297 1247053 459211]
[51247 636122 261997 ... 755708 521831 792985]
[993794 286242 565969 ... 659036 765881 641545]
[569821 769165 680852 ... 735125 779739 959945]
[605713 1404682 976885 ... 740845 944513 477987]
[97579 1294435 467748 ... 45698 1303553 756960]
[664463 29629 719953 ... 1427001 529414 1419260]
[1287116 636249 427210 ... 1253818 1075314 715063]
[757228 1274626 334159 ... 279499 1424781 1428208]
[175747 930739 953735 ... 1161898 1116252 959795]
[983968 488997 46597 ... 1185829 930460 1208424]
[1361592 69999 591252 ... 1163116 49587 951567]
[61837 975933 864424 ... 1397180 1001041 946954]
[974121 479586 255236 ... 1308838 9128 591585]
[716991 949564 717526 ... 490460 886779 1365359]
[493381 150331 1034749 ... 1358965 672295 985807]
[283845 859449 1447358 ... 1258167 233250 436222]
[982817 356944 1454024 ... 330816 1425310 181151]
[476331 978760 897728 ... 45593 499549 719541]
[677586 738497 1175960 ... 1066365 731885 378190]
[872939 857057 181534 ... 872671 1663 1402775]
[45960 418015 50944 ... 1328696 836912 368057]
[443113 1183829 885177 ... 974229 614334 1415379]
[355546 845810 492055 ... 347048 112993 152460]
[487304 676929 1044852 ... 504399 208439 1213199]
[1428359 671319 643070 ... 1409111 357178 82470]
[32192 860616 167497 ... 655095 1031646 1002204]
[766056 1146619 917254 ... 939197 676145 1310461]
[1238937 1098531 774903 ... 1053023 72965 418299]
[1242303 505996 1300313 ... 696197 1304215 1095193]
[89426 184651 787849 ... 270194 1046648 745780]
[1160325 834564 353747 ... 1211148 510153 876046]
[452027 811146 805860 ... 825108 384110 322750]
[123501 411749 842365 ... 432345 195655 338327]
[1342797 361052 453262 ... 577745 290144 1302703]
[1368190 1277343 397451 ... 719887 66039 187308]
[656832 803070 281982 ... 763158 725581 227831]
[1156313 961013 313496 ... 216963 669786 220843]
[1380000 641153 952141 ... 587820 1351046 619124]
[494407 312887 820859 ... 1447779 1411868 276948]
[1108418 989178 389155 ... 700877 1230786 442667]
[87023 249857 1161162 ... 1100196 250857 798845]
[677355 165861 1048584 ... 884577 211382 688622]
[1010041 438591 843989 ... 710285 253748 189112]
[579724 1299307 414920 ... 171022 1310999 1090755]
[1156825 1336039 162910 ... 858178 336273 1066958]
[824044 1373202 72400 ... 1394963 1004906 1087592]
[1216310 798048 1078414 ... 261018 1096595 802039]
[224611 316911 662692 ... 554168 699407 510108]
[238799 322937 579349 ... 845988 1024206 399991]
[394539 485123 1096147 ... 1073816 1415213 908968]
[598277 1035679 1311569 ... 268564 178296 615053]
[414048 886174 1096776 ... 1404711 679012 1249364]
[1133408 1195699 136425 ... 798810 1022772 987490]
[66851 885814 294697 ... 1073227 620932 630341]
[975644 486284 1303351 ... 1394504 339615 1092543]
[926371 237732 580636 ... 899039 695894 105346]
[47709 560631 794034 ... 1103234 721924 1151959]
[95622 571195 648057 ... 1349376 349776 937730]
[1082161 1398378 360932 ... 34291 1013879 623918]

```
[ 352656 1282263 79299 ... 1323511 259815 1110811]
[1215300 998254 341342 ... 1099292 1215581 595246]
[ 443669 928346 999042 ... 1010714 1071909 190991]
[ 953301 773781 346109 ... 1453316 265343 1020354]
[ 410902 1289046 1429213 ... 284940 1269382 1184667]
[ 427767 1418616 454440 ... 462357 474576 671479]
[ 645545 804430 1324636 ... 428144 332429 762827]
[1064735 1251350 562385 ... 688578 1371566 414722]
[ 937302 1142144 570637 ... 844122 660801 1019414]
[1203832 1402864 646885 ... 908482 481109 1388782]
[622465 283471 706428 ... 202711 133764 220826]
[1175705 1208527 1193103 ... 296395 176641 1078478]
[ 774639 397155 812513 ... 1138326 1009635 1378101]
[742162 10341 897202 ... 984904 800333 799714]
[ 245072 1399541 856299 ... 200808 716042 1325927]
[ 873308 974438 452058 ... 1264068 1058405 346096]
[ 255207 1327077 587955 ... 91385 323130 639814]
[393722 403170 297684 ... 685136 650412 115588]
[ 586775 680916 1075386 ... 184380 1455744 318247]
[ 529183 4582 205242 ... 1002904 710015 157413]
[1361544 353017 321029 ... 425699 1334477 474951]
[ 770879 1390030 782581 ... 70825 590291 465473]
[1455195 424385 187844 ... 221484 341448 175547]
[ 366292 980816 260365 ... 1247176 19642 352156]
[1167245 111012 41229 ... 638733 778621 1264197]
[ 312166 1293266 253087 ... 1274547 222766 995451]
[ 372170 381616 133265 ... 323696 401585 1019053]
```

```
In [ ]: accuracy = (prediction == y_test).sum() / prediction.shape[0]
accuracy
```

```
Out[ ]: 0.985858666204043
```

```
In [ ]: from sklearn.ensemble import AdaBoostClassifier
from sklearn.metrics import accuracy_score

# Create adaboost classifier object
abc = AdaBoostClassifier(n_estimators=100, learning_rate=1, random_state=42)

# Train Adaboost Classifier
model1 = abc.fit(X_train, y_train)

#Predict the response for test dataset
y_pred_abc = model1.predict(X_test)
```

```
In [ ]: accuracy = (y_pred_abc == y_test).sum() / y_pred_abc.shape[0]
accuracy
```

```
Out[ ]: 0.9858230094442414
```

```
In [ ]:
```